



Generative AI, Academic Integrity, and Introductory American Government: Can We Rebuild What You Destroy?

Charles C. Turner, Audrey B. Thao & Ginny Adorador

To cite this article: Charles C. Turner, Audrey B. Thao & Ginny Adorador (24 Jan 2026): Generative AI, Academic Integrity, and Introductory American Government: Can We Rebuild What You Destroy?, Journal of Political Science Education, DOI: [10.1080/15512169.2026.2619831](https://doi.org/10.1080/15512169.2026.2619831)

To link to this article: <https://doi.org/10.1080/15512169.2026.2619831>

 View supplementary material 

 Published online: 24 Jan 2026.

 Submit your article to this journal 

 Article views: 205

 View related articles 

 View Crossmark data 

 Citing articles: 1 View citing articles 



Generative AI, Academic Integrity, and Introductory American Government: Can We Rebuild What You Destroy?

Charles C. Turner, Audrey B. Thao and Ginny Adorador

California State University, Chico

ABSTRACT

Exponential growth in the capacities of Generative Artificial Intelligence (GAI) means college instructors face the challenge of being forward thinking while also providing opportunities that guarantee student learning. GAI is a rapidly changing technology and political science instructors need to know about attempted approaches and their relative success. Based on knowledge of student adaptation to technology, student desires for efficient and successful outcomes, and instructor desire for student learning, this article asks: can political science instructors develop assignments that reduce deceptive and inappropriate use of GAI, while encouraging both appropriate use of that technology and student learning? To answer this question empirically, we examine student responses to 16 low stakes writing assignments across two semesters, with prompts altered in the second semester to adapt for GAI. After coding responses based on likelihood of inappropriate GAI use, means are compared to identify more and less successful adaptations. This research should aid political science educators seeking positive uses of emerging technology and point out pitfalls to avoid. It will also benefit future students by encouraging faculty to guide students toward responsible uses of GAI and the development of their own thinking skills and political views rather than simple “cut and paste” techniques to complete assignments.

ARTICLE HISTORY

Received 8 July 2025
Accepted 30 December 2025

KEYWORDS

Generative artificial intelligence; AI; ChatGPT; pedagogy; academic integrity; introduction to American Government

Introduction

We have all been there. You start reading a student writing submission and something just seems...*off*. It doesn't sound like student writing. It doesn't even really sound like *human* writing at all. The response to the prompt you have written is not necessarily incorrect, but it feels hollow; it lacks depth. Then it dawns on you. The student did not write this. They copied your prompt into a Large Language Model (LLM) and then simply pasted the response into your assignment. You asked a good question, they provided an acceptable answer, and yet no learning took place. Panic sets in. What if other students did this as well? In fact, what if *every* student did this? Is higher education as we know it over? Is it time to throw in the towel?

This article aims to provide evidence that the answer to these last two questions is no. With some reflection on both the rapid development and ultimate limitations of LLMs and Generative Artificial Intelligence (GAI), political science instructors may be able to adapt assignments in ways that take advantage of the best features of GAI while still encouraging students to think critically about their role in a democratic society. To that end, we begin with an overview of current knowledge on student behavior and the development of GAI, then describe an empirical test of the hypothesis that assignments explicitly recognizing the world of GAI can reduce *unwanted* GAI use, and conclude with a discussion of results and suggestions for both the classroom and for further research. In other words, we argue that stating clear intentions about GAI and modeling its responsible use within the specific assignment instructions may be a productive path forward for student learning.

GAI presents a generational-level challenge to student learning in higher education. It raises many important questions that academics must begin to address both theoretically and empirically. Some of these are broad, theoretical questions about access and acceptability that will largely be answered at the institutional or societal level in the years to come. While such questions are undeniably important, they lie beyond the control of the typical college faculty member. On the other hand, some issues are—at least in part—within faculty control. For example, specific course assignments are typically the domain of the course instructor. Given the rapid development of GAI, faculty may need to change how they think about assignment development. We propose that one of the most vital questions individual faculty members can address is: “what can I do as an instructor to promote student thinking and discourage uses of GAI that do not promote student thinking?”

What we know about academic integrity, generative AI, and student behavior

Thanks to an expansion of publication venues for pedagogy-based research, political scientists have expanded the discipline’s understanding of a wide range of teaching and learning issues over the past generation. Today, academic integrity, student learning, and the use of technology all receive thoughtful empirical and theoretical scrutiny. Academics have been concerned about honesty and integrity generally, and the relationship between academic integrity and technology specifically, for at least a generation (see, e.g., Turner 2005; Lang 2013; Bertram Gallant and Rettinger 2022). In late 2022, the advent of readily available GAI tools introduced an important new variable into this broader conversation. Not surprisingly, university instructors around the globe have varying concerns about the impact of GAI on traditional conceptions of academic integrity (Lee et al. 2024; Yusuf, Pervin, and Román-González 2024; McMurtrie 2025).

At this point, political science instructors have likely all seen the horror stories: GAI can pass legal and medical exams, GAI hallucinates citations, GAI writing is indistinguishable from human writing. Alongside these concerns, market projections suggest that the number of GAI products focused on academia—numbering over one hundred by mid-2024—will continue to grow rapidly (Baytas and Ruediger 2024). Many of the nation’s largest higher education systems have taken notice and begun planning strategies to address this development (see, for example, California Community Colleges 2024; Palmer 2025).

Closer to home, we now know that GAI use is prevalent in collegiate political science courses (Cahill and McCabe 2024) and can be used to write competent essays (Dilling and Owen 2025), compile solid Model United Nations position papers (De Maio et al. 2025), and illustrate both strong and weak points of literature reviews (Tuncel 2025). The discipline as a whole has taken note of this technology's ability to transform how we teach (APSA Presidential Task Force 2024) and its potential for undermining our values, including social and environmental ones (Kehlenbach 2025). Perhaps most troubling, we know that GAI is good at providing decent answers to the types of questions we have historically asked, especially for entry-level topics, like those typically addressed in Introduction to American Government courses (Spindel and Ackerman 2025).

Though some institutions have attempted broad directives either restricting or condoning the use of various GAI tools, for the most part faculty have been left to their own devices to develop GAI-usage policies for their courses (Wu and Wu 2024). While some political science faculty have recommended policies banning GAI use entirely in the introductory American government classroom (Spindel and Ackerman 2025, 9) or have attempted to employ GAI-detection software (Ardoin and Hicks 2024), such approaches are at best naïve and at worst, biased. Studies have shown that students are using GAI whether we want them to or not, and often acknowledge this when surveyed (Cahill and McCabe 2024).¹ An extensive survey of Generation Z has also demonstrated that today's typical college-aged student uses GAI in their daily lives, even though many remain anxious and skeptical regarding this technology's ability to help their critical thinking skills (Gallup 2025). Moreover, there is compelling evidence that GAI-detection tools are flawed and may perpetuate bias as well (D'Agostino 2023; McMurtrie 2023). One scholar likened the detection-based approach to "an academic arms race" that instructors are unlikely to win (Kassorla 2023), while another notes that the advent of GAI is "so big that it literally changes the very premise of learning as we know it today" (Andriole 2024).

Another reason to avoid an all-out ban on the use of GAI is that this would throw out the good with the bad. Scholars have noted the value in pushing some of the more mundane aspects of writing and factual learning onto GAI in order to free up more instruction time to focus on understanding complex relationships and the development of effective research questions (Kassorla 2023). Others have pointed to GAI's value for enhancing accessibility and inclusion (Gonzalez et al. 2024) and have pointed to its worth as a tool to help us achieve outcomes we already value (Michels 2024; Kendrick 2025).² Finally, there is much speculation that skills with GAI tools will become important to the types of careers our students will pursue and "students will need to be taught the difference between what bots can do and what is expected of them as learners and professionals" (Michels 2024, 512). In short, there are too many potential benefits to ignore GAI entirely.

While findings such as those reported in the paragraphs above may lead some to despair that we can never assign writing again, thankfully, more reasonable alternatives are beginning to emerge. Some, for example, have leaned into GAI's limitations in developing opinions and original thoughts, as opposed to reporting facts (see Spindel and Ackerman 2025, 3). Others have found that GAI has strengths in question and assignment development and can be employed as an ally in developing better assignment prompts (Onal and Kulavuz-Onal 2024).

Many—perhaps nearly all—traditional writing prompts can no longer ensure student learning, nor can they help distinguish between students and machines. That being said, revising writing prompts may provide a solution that enhances the likelihood of student learning, and survey data suggest this is the most typical approach among political science faculty (Ardoin and Hicks 2024). To be sure, other solutions—such as the use of in-class writing prompts—can likely circumvent the use of GAI entirely, while maintaining critical thinking goals (Butcher 2022). Such solutions may have other drawbacks, however, that the instructor wishes to avoid. For example, they reduce class time that can be used for other objectives and they may create an unsustainable grading burden for large format courses like the introductory American government course.

So, if writing prompt revision is among the most viable approaches, how should prompts be revised? Since explicitly banning the use of GAI tools in the introductory American government classroom is unlikely to be effective, and would miss out on the genuinely useful features of GAI, a more effective approach might be to develop writing prompts that reduce the attractiveness of inappropriate GAI use (Ardoin and Hicks 2024) and/or “actively incorporate AI into the classroom” (Rivetti, Banerjee, and O’Mullane 2025; Spindel and Ackerman 2025, 2). What follows is an empirical examination of an effort to revise writing prompts for an age of accessible GAI.

Research question

When considering adoption and incorporation of GAI technology, Hodges and Ocak (2023) recommend conducting a technoethical audit.³ One of the components of such an audit is to ask questions like “in what ways does this technology encourage and discourage learning?” (Hodges and Ocak 2023).⁴ This question can be, at least in part, an empirical one. Thus, based on knowledge of student adaptation to technology, student desires for efficient and successful outcomes, and instructor desire for student learning, we ask the question: can instructors develop assignments that will reduce the inappropriate use of GAI, while encouraging the appropriate use of that technology? How might such adaptations affect student learning and engagement?

In order to best answer this question, and to avoid the problems of indirect observation of academic integrity issues such as surveys (see Wu and Wu 2024), this project employs a direct, empirical approach. We examine student responses to a series of low stakes writing assignments.⁵ These “chapter reflections” each provide a prompt that asks a student to spend two or three sentences reflecting on a concept from the relevant chapter (e.g., “what is a civil right or civil liberty that you would like to see the government do a better job of protecting?”). When these prompts were initially conceived (by one of the authors), the intention was to create a nexus between course material and the life of the student as a participant in a democratic society. Unfortunately, by the end of 2023, it was all too clear that GAI tools were well-equipped to provide adequate responses to these prompts, circumventing the need for students to reflect on the connection between American government and their own lives.

This reality led the lead author to experiment with revisions to these reflection prompts for the Fall 2024 semester. In particular, these revisions led to four types of

changes for comparison: no change from 2023 to 2024 (a baseline/control group), change to explicitly forbid GAI, change to discourage the use of GAI by focusing on the respondent's personal stakes in the issue, and changes that explicitly incorporate GAI into the assignment. With an eye toward the finding that “instructors are supportive of AI when it is used in an assistive as opposed to a generative manner” (Wu and Wu 2024, 605), this last category of revised prompts explicitly asks students to employ GAI tools as a first step in a larger endeavor.

Methods

Using data from two otherwise identical sections of an introductory American Government course, we employed a post-test only design with, effectively, a control group and a treatment group. Both courses were taught in a fall semester (one in 2023 and one in 2024) by the same instructor (the lead author) in an online, asynchronous format. Both courses were assigned a series of 16 weekly “chapter reflection” assignments to correspond with the 16 chapters in the course's textbook.⁶

The treatment itself was the revision to the assignment prompts discussed in the previous section (see [Supplementary Appendix A](#)). While a majority of the chapter prompts were reworded to explicitly incorporate or condone the usage of GAI tools, two chapter prompts were changed to make the use of AI less appealing by emphasizing the respondent's personal experience (chapters 10 and 15), one was changed to explicitly ban the use of GAI (chapter 1), and two were kept identical to the control group (chapters 4 and 8). In all cases, an effort was made to develop prompts that gave the student a stake in the answer by considering, for example, how an issue affects them personally, rather than simply ask the student to apply a broad concept generically.⁷ An additional element of the treatment was a shift in the syllabus statement on GAI use—from a prohibition approach in 2023 to a conditional use with transparency approach in 2024 (see [Supplementary Appendix B](#)). Transparency—in this case, the explicit acknowledgement of GAI use—has been identified by scholars as a best practice for AI policies (Lee et al. 2024).

Student responses to each of the 16 chapter reflections for both the control group (Fall 2023) and the treatment group (Fall 2024) were downloaded from the course Learning Management System (LMS) and de-identified into Excel files.⁸ Two of the coauthors—graduate research assistants with experience teaching the introductory American government course—then embarked on the time-consuming process of hand-coding each individual response as either “*unlikely* unacknowledged or inappropriate GAI use” (0), “*possible* unacknowledged or inappropriate GAI use” (1), or “*likely* unacknowledged or inappropriate GAI use” (2). Though, as noted above, detecting AI use is often a vain effort at the individual level, having multiple coders look at a larger universe of cases with an eye toward patterns at the class level may be more fruitful. The descriptive statistics in [Table 1](#) reflect an average of the two coders' scores.⁹ While there were ~120 students in each course, the total number of cases for each assignment ranges from 92 to 115, for a total of 3,415 coded observations across the two courses.

As an example of a response coded 0 (unlikely use), consider this student submission for the Fall 2024 Chapter 15 prompt:

Table 1. Reflection assignment data before and after GAI revision.

		Mean Character Count 2023	Mean Character Count 2024	Mean GAI Score 2023	Mean GAI Score 2024	Change in GAI Score	Significance (2-tailed)
Type of revision	Chapter						
Kept identical	4	473	443	.000	.026	.026	.062
		<i>N</i> =110	<i>N</i> =115	<i>N</i> =110	<i>N</i> =115		
		472	464	.000	.014	.014	.201
Ban GAI	1	<i>N</i> =100	<i>N</i> =110	<i>N</i> =100	<i>N</i> =110		
		538	542	.013	.004	−.009	.406
		<i>N</i> =113	<i>N</i> =115	<i>N</i> =113	<i>N</i> =115		
Made GAI less appealing	10	527	547	.019	.031	.012	.516
		<i>N</i> =105	<i>N</i> =112	<i>N</i> =105	<i>N</i> =112		
		516	617	.213	.218	.005	.946
Incorporate GAI	15	<i>N</i> =101	<i>N</i> =108	<i>N</i> =101	<i>N</i> =108		
		696	685	.023	.035	.012	.558
		<i>N</i> =108	<i>N</i> =114	<i>N</i> =108	<i>N</i> =114		
	3	535	571	.028	.065	.037	.170
		<i>N</i> =106	<i>N</i> =115	<i>N</i> =106	<i>N</i> =115		
		702	753	.109	.091	−.018	.692
	5	<i>N</i> =96	<i>N</i> =110	<i>N</i> =96	<i>N</i> =110		
		581	704	.014	.057	.043	.086
		<i>N</i> =109	<i>N</i> =115	<i>N</i> =109	<i>N</i> =115		
	7	545	702	.032	.152	.120	.005**
		<i>N</i> =95	<i>N</i> =112	<i>N</i> =95	<i>N</i> =112		
		592	812	.045	.079	.034	.307
	9	<i>N</i> =110	<i>N</i> =107	<i>N</i> =110	<i>N</i> =107		
		541	522	.039	.023	−.016	.560
		<i>N</i> =103	<i>N</i> =110	<i>N</i> =103	<i>N</i> =110		
	12	628	679	.087	.167	.080	.140
		<i>N</i> =92	<i>N</i> =105	<i>N</i> =92	<i>N</i> =105		
		429	526	.287	.138	−.150	.001**
	13	<i>N</i> =101	<i>N</i> =109	<i>N</i> =101	<i>N</i> =109		
		551	538	.165	.136	−.030	.595
		<i>N</i> =106	<i>N</i> =107	<i>N</i> =106	<i>N</i> =107		
	16	607	670	.175	.475	.299	.000**
		<i>N</i> =97	<i>N</i> =99	<i>N</i> =97	<i>N</i> =99		
		558	611	.078	.107	.029	.013*
total average		<i>N</i> =1652	<i>N</i> =1763	<i>N</i> =1652	<i>N</i> =1763		

p* < .05; *p* < .01.

I think that tax reforms would be extremely effective in assisting the poor in America. The rich elite has more money than I believe they can physically deal with in their lifetimes, and the idea of hoarding it while other people starve I believe is morally abhorrent. Nobody should be a billionaire. Period. Give us food or eat the rich.

While GAI could certainly have come up with that (or nearly any) response, it seems unlikely. The tone and amount of detail would have required multiple iterations with the GAI tool. Many responses coded 0 also contained specific information about the student’s community, school, or family, which GAI tools would be less likely to initiate. The assumption was that students looking for a shortcut would not employ GAI in a manner that would take much longer than simply writing down their own thoughts. On the other hand, consider this response to the same prompt, coded 2 (likely use):

Certainly! Here’s a cleaned-up version of your paragraph: I believe that the children of today are the future of tomorrow. Investing in education and vocational training programs at a young age allows individuals to expand their knowledge and develop their unique skills. It’s important to recognize that while not everyone excels at everything, each person brings something valuable to the table. By providing young people with opportunities to

enhance their intellectual abilities, we give them the gift of knowledge (something that can never be taken away). They can then use this knowledge to their advantage, pursuing higher-paying professions in fields they are passionate about.

Not only is this response vague and platitudinous, but the respondent also failed to remove the first part of the GAI tool's response when copying and pasting this answer. GAI tools like ChatGPT are also fond of bullet lists and bold text. While not dispositive, answers with such characteristics invited further scrutiny. Some responses also contained GAI hallucinations—naming, for example, fictional legislators when students were asked to name their congressional representative.

Many student responses fell somewhere between these two extremes. These “in-between” cases—coded 1—were responses that, while possibly student developed, contained language that led the coders to suspect GAI use. For example, responses sometimes used language that suggested their answer came from somewhere other than the textbook. The Chapter 16 prompt asks students, in part: “What principles should guide American foreign policy in the future?” While most students provided answers that described some aspect of realism, idealism, or isolationism (terms from the chapter), a few students used specific terminology that, while perfectly reasonable as an answer, did not appear in their book (e.g., “strengthening alliances and partnerships”). A student might come to this conclusion based on their general knowledge, or from another class they have taken, but this is also the type of answer a GAI tool might provide when asked the same question.

Another strategy that led to responses being coded 1 was a shift in tone between different parts of a multi-point prompt. For example, the Chapter 8 prompt asks students whether they are registered to vote and why, as well as whether they plan to vote in an upcoming election and why. This led to some cases being coded 1 when the answer to one part of the prompt was deeply personal (e.g., wanting to vote against a politician who threatens aspects of the student's identity) while the answer to the other part was highly generic (e.g., supporting democracy). Since these shifts in tone raised suspicions, but could result from causes other than GAI use, they received this *possible* use code.

Limitations of this design include the fact that the two course sections were taught in different years, that one of those years was a presidential election year, and that Generative AI tools became more advanced and more widely available by the second year.¹⁰ In particular, the shift from ChatGPT 3.5 to ChatGPT 4 that occurred between the two years meant that more students had access to a free version of the more advanced tool during the second year.¹¹ On one hand, it would have been ideal to have selected both groups from the same semester, but teaching schedules did not allow for this. Moreover, there would have been something disingenuous about employing flawed assignments once the instructor knew of (or at least suspected) their limitations. Additionally, syllabus language regarding the use of AI differed between semesters. These differences can be seen in [Supplementary Appendix B](#).

As a final, and significant, note of caution, we are aware that the exercise of trying to identify human-generated *versus* AI-generated answers has often resulted in outcomes little better than chance (Wu and Wu 2024), and this problem may have been exacerbated by the advent of ChatGPT 4 and, now, ChatGPT 5. We have attempted to

counteract this difficulty in four ways. First, averaging the results of two coders lessens the impact of errors. Second, reading field-specific responses has been shown to be more effective than those written on generic topics (Wu and Wu 2024). Third, examining short writings, as opposed to longer essays should make the AI speech patterns more obvious.¹² Finally, the relevant outcomes are comparative rather than absolute in nature. In other words, the key finding is not whether a particular response was AI-generated, but whether an entire set of responses was, on average, more affected by unsanctioned GAI use than another set. Just as some students have avoided detection of more traditional plagiarism in the past, some will avoid detection of inappropriate GAI use. That said, the method employed here—looking at large numbers of small assignments—helps identify the broader pattern.

Findings

The findings displayed in [Table 1](#) warrant a few observations. Despite the discouraging nature of the overall findings, we begin by noting some positives. One unanticipated bright spot was the change in average length of student responses, which grew, on average, from 558 characters to 611 characters—a 9.5% increase—between Fall 2023 and Fall 2024. Though conjecture, we suspect that the growth in response length was a result of allowing/encouraging the use of AI tools (as this increase did not occur in the chapters where the writing prompt either remained unchanged or explicitly forbade the use of AI). We mention this not only as an encouraging outcome, but also as a word of caution to instructors: student submissions may increase in length with the incorporation of GAI. So, it may be useful to either institute a word limit or prepare for more reading.

A second optimistic finding is that the vast majority of students were probably not using GAI inappropriately. Across the total number of observations, over 90% were coded as unlikely to have used GAI inappropriately. Despite the justified academic-integrity concerns of educators, most students in this study used GAI tools either responsibly or not at all. This outcome is consistent with our assumption that students are unlikely to risk negative consequences for inappropriate GAI use in brief, low stakes assignments. Moreover, while harder to identify with the data provided here, the design efforts of the instructor might well make a difference. We note that, while not statistically significant, assignment modification led to a decrease in the average GAI score for five of fourteen assignments. In the case of the two unmodified assignments, averages rose. Moreover, the three chapters with significant results all involved assignment revision. In other words, what we do as instructors matters; we just need to continue researching how to make it matter in the way we intend.

Breaking the differences down by strategy merits consideration as well. Here, we will focus specifically on assignments that were revised to incorporate GAI as a tool before asking students to submit their own reflections (see [Supplementary Appendix A](#)). Three assignments making such revisions produced statistically significant results; in one case, inappropriate GAI use decreased and in two cases it increased. The Chapter 7 prompt was changed from asking students a general question about their involvement with interest groups that referenced the textbook to one that instructed students to make use of the internet or GAI tools to generate a list of interest groups

working in a field of their personal interest and then discuss which one they might want to join in the future. The Chapter 16 prompt shifted from asking students a general question about the history of American foreign policy to one that asks students to use GAI to generate a list of current world conflicts before selecting one to analyze in more detail. In both of these cases, inappropriate use of GAI increased. The Chapter 13 prompt shifted from a factual question about the five stages of the policy process to a question that asks students to consider which part of that process they would most like to work on in their career. In this case, inappropriate GAI use decreased significantly.

So, why did one of these prompt revisions succeed in its objective while the other two made the problem worse? While this question merits further investigation, two possibilities are worth consideration. One, the two prompts that made the situation worse both explicitly asked the respondent to use a GAI tool but then failed to explicitly state that students should then use that information to craft an answer using their own words. Even though this distinction is clearly spelled out in the course syllabus and in other assignments, it is possible that some students, after being instructed to enter the world of GAI, took license with the lack of a chapter-specific prohibition. Two, the successful prompt may have benefited from modeling the appropriate use of GAI tools. It twice mentions appropriate ways to employ internet/GAI tools before emphasizing that the response needs to reflect the student's own personal interests and be written in their own words. Perhaps the key takeaway here is that there is value in being as explicit as possible, as often as necessary, even when the rules seem obvious to the instructor.¹³ Overall, the prompt revisions in [Supplementary Appendix A](#) did not prevent, or even reduce inappropriate GAI use, but they may serve as a first step toward better practices.

Despite the optimism of the previous paragraphs, the findings also indicate need for concern. First, and most obviously, efforts to steer students away from inappropriate use of GAI were largely unsuccessful. Overall, there was a statistically significant *increase* in inappropriate use between Fall 2023 and Fall 2024. Across the 16 chapters, only five indicated a decrease in GAI use and in only one case was that decrease statistically significant. In total, about 8% of 2023 assignments were coded as at least "possible" GAI violations, a figure that rose to nearly 11% in 2024. In other words, despite a shift from an absolute ban on GAI to a more nuanced approach, and despite many efforts to model responsible uses of GAI in assignment prompts, the number of students choosing GAI as an unsanctioned shortcut increased.

Second, no particular approach to assignment revision was a clear success. There were no statistically significant differences in any of the chapters in which assignments were kept identical, in which GAI was explicitly banned, or in which the prompt language was personalized in an effort to make GAI use less appealing. Efforts to incorporate GAI use into the assignments demonstrated the most potential for changing student behavior, though not necessarily in the manner intended.

Finally, in addition to these quantifiable differences, the content of student responses displayed some qualitative differences which, though likely more subjective and open to interpretation, are worth noting. For one, though the student responses had been de-identified, both coders concluded that it was likely that a small number of students were using GAI on every assignment. A pattern of unusual and memorable answers

structures (e.g., beginning each chapter reflection with the same wording) led to this conclusion, raising the possibility that some subset of students may have decided that neither the risk of punishment nor the reward of learning was large enough to approach these assignments with integrity.

We also observe that, though GAI is becoming much more prevalent, all students are not users. Several students in the fall 2024 cohort mentioned these assignments as their first time using GAI, at least in an academic setting. Such statements, combined with the low numbers of inappropriate usage described above, should serve as a reminder that students need instruction in using these tools. While nearly all students have heard of GAI tools by this point, and some are likely to be using them in unacceptable ways, many may be eager to see demonstrations of, and have opportunities to practice, responsible use. In sum, the data suggest that, while fine tuning is needed, there is potentially more to be gained by engaging with GAI in the American government classroom than by simply ignoring or banning it.

Discussion

The findings above suggest a few key takeaways. First, there is likely no GAI-proof assignment; some subset of students will continue to look for shortcuts (like GAI), even when instructors go to great lengths to personalize assignments. Consider the Chapter 4 prompt, which the instructor did not change between 2023 and 2024 due to its track record of producing thoughtful responses: “Take this political ideology questionnaire: <https://www.theadvocates.org/quiz/> What result did you get? Is this consistent with your view of yourself? Why or why not? Your answer should be a few sentences.” Even this assignment that requires no prior knowledge and only benefits the student when done personally has seen a rise in unsanctioned GAI usage.¹⁴

Second, developing assignments that discourage GAI use might lead one to develop assignments that are more focused on student opinion and less about “critical thinking” as traditionally understood. Indeed, it is likely that the revisions employed here have done exactly that. Unfortunately, introductory level objective critical analysis is an area where GAI excels. While faculty should certainly continue to develop assignments that develop this important skill, they may need to do so in formats that can more easily avoid GAI use, such as in-class writings. Thinking through one’s own political views in an informed manner is also an important skill for introductory political science students, and may be well-suited to the types of low-stakes assignments described here. Figuring out how to use AI tools to inform and help develop political opinions may become a very useful tool for building civic skills.

Third, inappropriate GAI use will probably continue to rise. As more AI natives enter college, and as GAI tools continue to increase their capabilities and expand their reach, their use—as both an appropriate tool and as an inappropriate temptation—will surely become ubiquitous. The storm of Generative AI has arrived in higher education. We cannot stop it, but we may be able to stem and redirect its more deleterious effects. The more explicit we can be about how and when GAI can be educational, the better.

Since ignoring the existence of GAI and simply banning its use are strategies unlikely to produce desired outcomes, political scientists should embrace options that guide students in a positive direction by very explicitly indicating when and when not to use GAI tools and by encouraging students to use GAI in ways that will help them get started with larger projects that require human-generated analytical thinking and political opinion formation. Indeed, since the introductory American government course often has the development of informed ideas and opinions as one of its main goals, the rise of GAI is also a perfect moment to encourage and challenge instructors to develop assignments that point students toward this goal. As this study demonstrates, it will take some trial and error to develop the right mix of assignments and incentives. This mix may differ in some ways by course, instructor, and institution, but if we can harness GAI as a tool for building confidence in the substantive content of American government, we may be encouraging development of confident, informed political opinions.

To be sure, Generative AI and other technologies will continue to advance, and some subset of college students will always be looking for a shortcut, but one conclusion of this research is that there is hope. There are ways to rebuild what we value, and what GAI, with one of its undesirable byproducts, has destroyed. The approach described here—explicit instructions and modeling within specific assignments—is just one such way and it is the authors' hope that instructors reading this will not despair as they gaze into the abyss of GAI, but will be inspired to develop new approaches for student learning.

We suspect that, in the long run, society will benefit from the presence of educated citizens who know how to responsibly differentiate between tasks they can automate and those that require human thought and reflection. It is our hope that this research will aid political science educators in the quest to identify positive and helpful uses of emerging technology and point out pitfalls that they would be best to avoid. Through further empirical studies of how various GAI strategies work in pedagogical practice, we can build a knowledge base of how GAI can aid learning in the American government classroom. We should also focus on what AI *cannot* do—it cannot maintain a healthy democracy. Only active, informed humans can do that. So, let us focus on developing more of those.

Notes

1. Though, Wu and Wu (2024, 602) note, “it is difficult to obtain reliable changes in LLM usage by students due to social desirability bias and sampling issues.”
2. Though, to be sure, social scientists also have concerns about the political biases that may be present in LLMs (Wu and Wu 2024; Motoki et al. 2025). Though such an investigation is beyond the scope of the present study, we are indebted to an anonymous reviewer for raising this concern and recommend further investigation into how such bias may impact political science students.
3. Hodges and Ocak (2023) describe a technoethical audit as “a critical evaluation of the chosen technology. Such an audit explores whether it is ethical to use the technology and what potentially unfavorable outcomes might arise from its use in schools.”
4. Another important question to consider is, “what motivates a student to use a GAI tool?” While we acknowledge the importance of this question and encourage others to pursue it, it is not one we consider here.

5. Taken as a whole, these assignments comprise 20% of a student's grade. Each individual reflection comprises only 1.25% of the semester grade. We should emphasize here that, while these assignments require students to write something, their main purpose is not to improve students' writing abilities, but rather to improve students' abilities to reflect and think critically about what they have read.
6. Students in both courses were assigned the online lab book version of *Introduction to American Government*, 11th ed. (Turner et al. 2021).
7. As Michels (2024, 517) notes: "having these low-stakes types of general writing on big themes and concepts would seem to invite copying and pasting from a bot. Instructors wanting to continue these kinds of assignments should consider asking more questions about real-life situation and the personal experiences of students. AI does not perform as well with these sorts of prompts."
8. Canvas was the LMS used for both course sections.
9. Prior to coding, the two coders were assigned reading on GAI detection strategies (see, for example, East Central College. 2024), engaged in norming exercises, and were encouraged to explore GAI tools to see the types of responses the chapter prompts might provide. The inter-coder reliability (Cohen's kappa) rating of .359 (*SE* .049) for 2023 and .438 (*SE* .041) for 2024 indicates fair to moderate agreement, likely suffering from the "Kappa Paradox" due to the large number of cases that were coded as 0 (see Schnell n.d.). The overall agreement rate of 93.9% for 2023 and 92.6% for 2024 suggests strong agreement. Ultimately, we are more concerned with AI-likelihood than AI-certainty, so the mean score across two coders is a reasonable measure.
10. It is also possible that there is a time-based effect within each semester: over the course of the semester, as stakes become higher and time becomes more precious (due to other obligations in school and work), students may be more willing to use GAI in inappropriate ways. Since this effect would apply to both semesters equally, we do not pursue that possibility here; we simply observe that it may be worthy of future study.
11. Though some assignment instructions referred to a specific tool, like ChatGPT or Copilot, these were used as examples since they were the tools students were likely to be the most familiar with. The course syllabus and class instruction addressed all GAI tools in a similar fashion and no effort was made to distinguish between them as more or less desirable/appropriate for the course.
12. Though, to be sure, since the coders were not familiar with the students' broader writing styles, they may have coded more conservatively than an instructor more familiar with the students.
13. We also note the value of providing students with a specific location to look for information. Chapters 1, 5, 11, and 13, for example, all explicitly point students to the text and inappropriate GAI usage decreased (though not necessarily significantly) in each of those chapters. Future research may want to consider whether simply reminding students of the connection to assigned course materials may be an effective deterrent.
14. The difference in means is significant at the .05 level in a one-tailed test.

Acknowledgments

With appreciation and acknowledgement to The Linda Lindas, whose lyrics always get one thinking.

Ethical approval

This research has institutional review board approval from the authors' institution: IRB-2024-112. Student consent was obtained for all quoted assignment responses.

Disclosure statement

No potential conflict of interest was reported by the author(s).

Funding

The authors wish to acknowledge support from the Office of Faculty Development at California State University, Chico. We also acknowledge research assistance funding from the Department of Political Science and Criminal Justice at California State University, Chico.

Notes on contributors

Charles C. Turner is a professor of Political Science at California State University, Chico, where he teaches courses on American political institutions and coordinates the Master of Public Administration program. He is a coauthor of *Introduction to American Government* and is a member of the California Faculty Association.

Audrey B. Thao is a Master's student studying Political Science at California State University, Chico after completing her Bachelor's in Political Science at University of California, Davis.

Ginny Adorador is a teaching associate in Political Science at California State University, Chico, where she teaches discussion sections for introductory American government and supports student learning through tutoring and supplemental instruction. She is completing her master's degree in political science, with an expected graduation date of May 2026. Her academic work has been recognized with multiple honors, including the Michele J. Shover Academic Excellence in Political Science Award.

References

- Andriole, Steve. 2024. "How Generative AI Owns Higher Education. Now What?" *Forbes*, March 18. <https://www.forbes.com/sites/steveandriole/2024/03/18/how-generative-ai-now-owns-higher-education--now-what/> (accessed October 17, 2024).
- APSA Presidential Task Force. 2024. *Rethinking Political Science Education*. March. <https://educate.apsanet.org/wp-content/uploads/APSA-Presidential-Task-Force-Report-on-Rethinking-Political-Science-Education-March-2024.pdf> (accessed January 21, 2025).
- Ardoin, Phillip J., and William D. Hicks. 2024. "Fear and Loathing: ChatGPT in the Political Science Classroom." *PS: Political Science & Politics* 57 (4):583–593. doi: 10.1017/S1049096524000131.
- Baytas, Claire, and Dylan Ruediger. 2024. "Generative AI in Higher Education: The Product Landscape." *Ithaka S+R*, March 7. <https://sr.ithaka.org/publications/generative-ai-in-higher-education/>.
- Bertram Gallant, Tricia, and David Rettinger. 2022. "An Introduction to 30 Years of Research on Academic Integrity." *Journal of College and Character* 23 (1):1–5. doi: 10.1080/2194587X.2021.2017975.
- Butcher, Charity. 2022. "Using In-Class Writing to Promote Critical Thinking and Application of Course Concepts." *Journal of Political Science Education* 18 (1):3–21. doi: 10.1080/15512169.2021.1990073.
- Cahill, Christine, and Katherine McCabe. 2024. "Context Matters: Understanding Student Usage, Skills, and Attitudes Toward AI to Inform Classroom Policies." *PS: Political Science & Politics* 57 (4):594–601. doi: 10.1017/S1049096524000155.
- California Community Colleges. 2024. "Generative AI and the Future of Teaching and Learning." Report to Board of Governors, July. <https://www.cccco.edu/-/media/CCCCO-Website/docs/report/generative-ai-and-the-future-of-teaching-and-learning-7-17-24-2-a11y.pdf>.
- D'Agostino, Susan. 2023. "Turnitin's AI Detector: Higher-Than-Expected False Positives." *Inside Higher Education*, June 1. <https://www.insidehighered.com/news/quick-takes/2023/06/01/turnitins-ai-detector-higher-expected-false-positives> (accessed January 24, 2025).
- De Maio, Jennifer L., Ismail Kabalaki, Shayan Moshtael, and Michael A. Tejax. 2025. "AI versus Students: A Study of the Capability of ChatGPT to Write Model United Nations Position Papers." *PS: Political Science & Politics* 58 (2):357–362. doi: 10.1017/S1049096524000799.

- Dilling, Matthias, and Leah Owen. 2025. "Designing Politics and IR Assessments in the Era of AI: An Empirical Investigation into ChatGPT's Output Across Bloom's Revised Taxonomy." *Journal of Political Science Education* 21 (2):290–309. doi: [10.1080/15512169.2024.2408767](https://doi.org/10.1080/15512169.2024.2408767).
- East Central College. 2024. "Detecting AI-Generated Text: 12 Things to Watch For." <https://www.eastcentral.edu/free/ai-faculty-resources/detecting-ai-generated-text/> (accessed January 21, 2025).
- Gallup, Inc. 2025. "Voices of Gen Z: How American Youth View and Use Artificial Intelligence." April 1. <https://www.gallup.com/analytics/658901/walton-family-foundation-gallup-voices-gen-american-youth.aspx> (accessed May 20, 2025).
- Gonzalez, Lorna, Kristi O'Neil-Gonzalez, Megan Eberhardt-Alstot, Michael McGarry, and Georgia Van Tyne. 2024. "Leveraging Generative AI for Inclusive Excellence in Higher Education." *EDUCAUSE Review*, August 15. <https://er.educause.edu/articles/2024/8/leveraging-generative-ai-for-inclusive-excellence-in-higher-education> (accessed October 17, 2024).
- Hodges, Charles, and Ceren Ocak. 2023. "Integrating Generative AI into Higher Education: Considerations." *EDUCAUSE Review*, August 30 (accessed October 17, 2024).
- Kassorla, Michelle. 2023. "Teaching with GAI in Mind." *EDUCAUSE Review*, December 14. <https://er.educause.edu/articles/2023/12/teaching-with-gai-in-mind> (accessed December 5, 2024).
- Kehlenbach, E. Stefan. 2025. "The Impact of Infrastructure: Considerations of Generative AI in the Classroom." *Journal of Political Science Education* 21 (2):267–276. doi: [10.1080/15512169.2024.2359439](https://doi.org/10.1080/15512169.2024.2359439).
- Kendrick, LaRae. 2025. "Our 'Six Traits' Roots and Rescue Revision Can Save Personalized Writing in Our New Age of AI." *LinkedIn*, April 29. <https://www.linkedin.com/pulse/our-six-traits-roots-rescue-revision-can-save-writing-larae-04wgc/> (accessed April 30, 2025).
- Lang, James. 2013. *Cheating Lessons: Learning from Academic Dishonesty*. Cambridge, MA: Harvard University Press.
- Lee, Daniel, Matthew Arnold, Amit Srivastava, Katrina Plastow, Peter Strelan, Florian Ploeckl, Dimitra Lekkas, and Edward Palmer. 2024. "The Impact of Generative AI on Higher Education Learning and Teaching: A Study of Educators' Perspectives." *Computers and Education: Artificial Intelligence* 6:100221. doi: [10.1016/j.caeai.2024.100221](https://doi.org/10.1016/j.caeai.2024.100221).
- McMurtrie, Beth. 2023. "How ChatGPT Could Help or Hurt Students with Disabilities." *The Chronicle of Higher Education*, May 26. <https://www.chronicle.com/article/how-chatgpt-could-help-or-hurt-students-with-disabilities> (accessed January 24, 2025).
- McMurtrie, Beth. 2025. "Is Your Online Student a Bot?" *The Chronicle of Higher Education*, May 1. <https://www.chronicle.com/article/is-your-online-student-a-bot> (accessed May 7, 2025).
- Michels, Steven. 2024. "Teaching (with) Artificial Intelligence: The Next Twenty Years." *Journal of Political Science Education* 20 (4):510–521. doi: [10.1080/15512169.2023.2266848](https://doi.org/10.1080/15512169.2023.2266848).
- Motoki, Fabio, Valdemar Pinho Neto, and Victor Rangel. 2025. "Assessing Political Bias and Value Misalignment in Generative Artificial Intelligence." *Journal of Economic Behavior & Organization* 234:106904. doi: [10.1016/j.jebo.2025.106904](https://doi.org/10.1016/j.jebo.2025.106904).
- Onal, Sinan, and Derya Kulavuz-Onal. 2024. "A Cross-Disciplinary Examination of the Instructional Uses of ChatGPT in Higher Education." *Journal of Educational Technology Systems* 52 (3):301–324. doi: [10.1177/00472395231196532](https://doi.org/10.1177/00472395231196532).
- Palmer, Kathryn. 2025. "Tech Giants Partner With Cal State System to Advance 'Equitable' AI Training." *Inside Higher Ed.*, February 5. <https://www.insidehighered.com/news/tech-innovation/artificial-intelligence/2025/02/05/cal-state-system-tech-giants-partner#> (accessed July 2, 2025).
- Rivetti, Paola, Rituparna Banerjee, and David O'Mullane. 2025. "Testing ChatGPT in International Relations Classrooms: Potentialities, Limitations, and What's Next." *PS: Political Science & Politics* 58 (1):169–176. doi: [10.1017/S1049096524000817](https://doi.org/10.1017/S1049096524000817).
- Schnell, Audrey. n.d. "What is Kappa and How Does it Measure Inter-Rater Reliability?" *The Analysis Factor*. <https://www.theanalysisfactor.com/kappa-measures-inter-rater-reliability/> (accessed October 31, 2024).
- Spindel, Jennifer, and James Ackerman. 2025. "Adapting to the AI Era: ChatGPT in the Classroom." *Journal of Political Science Education* 21 (2):310–326. doi: [10.1080/15512169.2024.2414017](https://doi.org/10.1080/15512169.2024.2414017).
- Tuncel, Ozlem. 2025. "Teaching Literature Reviews in the Age of Generative Artificial Intelligence (AI)." *Political Science Educator* 28 (2). <https://educate.apsanet.org/teaching-literature-reviews-in-the-age-of-generative-artificial-intelligence-ai>.

- Turner, Charles C. 2005. "A New Honesty for a New Game: Distinguishing Cheating From Learning in a Web-Based Testing Environment." *Journal of Political Science Education* 1 (2):163–174. doi: [10.1080/15512160590961810](https://doi.org/10.1080/15512160590961810).
- Turner, Charles C., D. Grier Stephenson Jr., Robert J. Bresler, Robert J. Friedrich, and Joseph J. Karlesky. 2021. *Introduction to American Government*. 11th ed. Redding, CA: BVT Publishing.
- Wu, Nicole, and Patrick Y. Wu. 2024. "Surveying the Impact of Generative Artificial Intelligence on Political Science Education." *PS: Political Science & Politics* 57 (4):602–609. doi: [10.1017/S1049096524000167](https://doi.org/10.1017/S1049096524000167).
- Yusuf, Abdullahi, Nasrin Pervin, and Marcos Román-González. 2024. "Generative AI and the Future of Higher Education: A Threat to Academic Integrity or Reformation? Evidence from Multicultural Perspectives." *International Journal of Educational Technology in Higher Education* 21 (1):1–29. doi: [10.1186/s41239-024-00453-6](https://doi.org/10.1186/s41239-024-00453-6).